
Diffusion scores for Markov data: exact inference on the chain and estimation of the transition kernel

Giovanni Mantovani
Department of Computing Sciences
Bocconi University
giovanni.mantovani@studbocconi.it

Internal note, work in progress. Circulated for discussion. The inference and estimation results are settled and reproducible; the generative comparison of §8 is complete for one innovation family with four more in progress, and its interpretation is the main open question. Section 9 states what is not yet supported. Open items for the advisors are listed separately in QUESTIONS_FOR_ADVISORS.md.

Abstract

Score-based generative models are trained by denoising regression, and the function they approximate — the posterior mean of the clean signal given its noised version — is unavailable in closed form for essentially every distribution of interest. This makes it hard to separate the error a network makes from the difficulty of the object it approximates. We study a setting where the posterior mean is computable and non-trivial: a first-order Markov chain on $\mathbb{R}^n$ with continuous, non-Gaussian transitions. Coordinatewise Ornstein–Uhlenbeck noising multiplies the prior by unary factors, so the posterior remains a chain, and sum-product inference on it is exact at the level of functional messages. We use this to obtain a controlled benchmark, and then remove the assumption that the transition kernel is known: it is estimated from noised sequences alone by maximum marginal likelihood, with an exact expectation step and a gradient supplied by Fisher’s identity, so the inference recursion is never differentiated through. On a Laplace chain the estimator attains a lower denoising error than trained neural denoisers at every data budget we tried, by roughly an order of magnitude against a fully connected network and a factor of four against a locality-respecting one. We then compare the same estimators through reverse diffusion, and the ranking does not carry over: the structured estimator remains the most accurate pointwise while the convolutional network reproduces the innovation law more faithfully. We report this dissociation as the principal empirical finding rather than resolving it, and we identify the information-theoretic asymmetry that plausibly explains it.

1 Introduction

Generative diffusion generates data $a \in \mathbb{R}^n$ from a distribution P_0 known only through samples, by noising and then reversing [Sohl-Dickstein et al., 2015, Ho et al., 2020, Song et al., 2021]. In the Ornstein–Uhlenbeck convention we follow throughout [Mézard, 2025], the forward process is

$$\frac{dx}{dt} = -x + \sqrt{2}\eta(t), \quad \text{so} \quad x(t) = e^{-t}a + \sqrt{\Delta_t}z(t), \quad \Delta_t = 1 - e^{-2t}, \quad (1)$$

with z standard Gaussian. Reversal is driven by the score $S(x, t) = \nabla_x \log P_t(x)$, and the central quantity is the posterior mean $\langle a \rangle_{x,t}$ of the clean signal.

Three distinct score objects. Confusing them is the source of a common misstatement, so we fix them at the outset.

The *population score* is $\nabla_x \log P_t(x)$ for the true P_0 . The *empirical score* replaces P_0 by the empirical measure $\hat{P}_0 = M^{-1} \sum_{\mu} \delta_{a^\mu}$ of the training set; after noising this is a Gaussian mixture centred on the training points. The *learned score* is what a network with finite capacity, finite data and a particular optimiser actually produces.

The population denoising objective $\mathbb{E} \|f(x, t) - a\|^2$ is minimised over all measurable f by the population posterior mean. The *empirical* objective, which is what is computed in practice, is minimised by the *empirical* posterior mean. Since exact reversal under the empirical score returns the empirical distribution, a model that attained the unrestricted optimum of its own training objective would reproduce its training set. That trained models generalise is therefore a statement about the restrictions — architecture, capacity, optimisation, regularisation — and not about the objective [Biroli et al., 2024, Kamb and Ganguli, 2024]. We do not claim this is a complete account of neural generalisation; we note it because it motivates wanting settings where the population score is available for comparison.

Why a Markov chain. Distributions with tractable population scores are scarce, which is why the analytical literature leans on Gaussians and Gaussian mixtures. Those are poor probes of structure: the score is linear, so any question about how a model exploits correlation has a degenerate answer. Hierarchical and discrete graphical models supply further tractable families [Sclocchi et al., 2025, Garnier-Brun et al., 2024, Mei and Wu, 2024], and the model here is a continuous, non-Gaussian member of the same class.

A first-order Markov chain gives a genuinely nonlinear score while remaining exactly computable, for a structural reason: coordinatewise noising contributes one unary factor per site, which reweights node potentials without creating edges between latent variables, so the posterior factor graph is the same chain as the prior’s.

Questions. We ask: (i) what does exact inference on the chain buy over the Gaussian approximation, and where in the noise schedule does it matter; (ii) can the transition kernel be estimated from noised data alone, and at what cost; (iii) how does the resulting estimator compare with trained denoisers, both pointwise and after reverse diffusion.

Contributions.

- An exact denoising benchmark for a continuous non-Gaussian correlated model (§3), with the discretisation error separated from the model error and measured (§4).
- An identification of what the single-Gaussian message closure computes: exactly the LMMSE estimator of the covariance-matched Gaussian model (§5).
- Estimation of the unknown kernel from noised sequences, with an exact expectation step and no differentiation through the recursion (§6).
- A denoising comparison against neural baselines (§7) and a generative comparison (§8) which do not agree, together with the information asymmetry that plausibly explains the disagreement.

2 Setting

Notation. We write n for the chain length, which is also the ambient dimension, and M for the number of observed sequences. (The source we follow for diffusion notation uses n for the number of samples and d for the dimension; we index by the chain because every graphical-model statement below is indexed by sites.) Sequences are $a = (a_1, \dots, a_n)$, and P_0 factorises as

$$P_0(a) = \mu(a_1) \prod_{i=2}^n W(a_i | a_{i-1}), \quad W(a' | a) = \varphi(a' - \rho a), \quad (2)$$

with φ an innovation density of mean zero and variance q , and $|\rho| < 1$. Choosing $q = 1 - \rho^2$ makes the chain stationary with unit marginal variance and $\text{Cov}(a_i, a_j) = \rho^{|i-j|}$ *independently of the shape*

of φ . This is what makes the family a controlled probe: Gaussian, Laplace, Student, uniform and bimodal innovations can be arranged to share a covariance exactly and differ only beyond second moments.

The score and the posterior mean. From (1), P_t is the Gaussian convolution

$$P_t(x) = \int da P_0(a) \frac{1}{(2\pi\Delta_t)^{n/2}} \exp\left[-\frac{\|x - e^{-t}a\|^2}{2\Delta_t}\right]. \quad (3)$$

Differentiating under the integral,

$$\nabla_x P_t(x) = \int da P_0(a) \frac{e^{-t}a - x}{\Delta_t} \frac{e^{-\|x - e^{-t}a\|^2/2\Delta_t}}{(2\pi\Delta_t)^{n/2}},$$

and dividing by $P_t(x)$, the Gaussian factor over $P_t(x)$ is exactly the posterior density $P(a | x, t)$, so

$$S(x, t) = \nabla_x \log P_t(x) = -\frac{x - e^{-t}\langle a \rangle_{x,t}}{\Delta_t}, \quad \langle a \rangle_{x,t} = \int da a P(a | x, t). \quad (4)$$

Differentiation under the integral is justified whenever P_0 has a finite first moment, by dominated convergence with the Gaussian kernel and its derivative. Equation (4) is a change of variables, not an approximation: computing the score and computing the posterior mean are the same problem.

The Bayes-optimal denoiser. For squared loss and any u that is a function of x alone,

$$\mathbb{E}[\|a - u\|^2 | x] = \mathbb{E}[\|a - \langle a \rangle_{x,t}\|^2 | x] + \|u - \langle a \rangle_{x,t}\|^2, \quad (5)$$

so $\langle a \rangle_{x,t}$ minimises the conditional risk, with an irreducible floor given by the first term. Optimality is with respect to squared error specifically; a different loss selects a different functional of the posterior.

3 Exact inference on the noised chain

Conditioning (2) on x gives

$$P(a | x, t) \propto \mu(a_1) \prod_{i=2}^n W(a_i | a_{i-1}) \prod_{i=1}^n \ell_t(x_i | a_i), \quad \ell_t(x_i | a_i) \propto \exp\left[-\frac{(x_i - e^{-t}a_i)^2}{2\Delta_t}\right]. \quad (6)$$

The likelihood factors are *unary*: because the forward noise is independent across coordinates, each observation touches one latent variable. They therefore reweight node potentials and add no edges, and the posterior factor graph is the prior's chain (Figure 1). A chain is a tree, so sum-product is exact [Pearl, 1988, Mézard and Montanari, 2009]: with

$$\begin{aligned} \vec{m}_i(a_i) &\propto \int da_{i-1} W(a_i | a_{i-1}) \ell_t(x_{i-1} | a_{i-1}) \vec{m}_{i-1}(a_{i-1}), \\ \overleftarrow{m}_i(a_i) &\propto \int da_{i+1} W(a_{i+1} | a_i) \ell_t(x_{i+1} | a_{i+1}) \overleftarrow{m}_{i+1}(a_{i+1}), \end{aligned} \quad (7)$$

initialised by $\vec{m}_1 = \mu$ and $\overleftarrow{m}_n \equiv 1$ — without which the initial law would enter no marginal — the single-site marginal is $P(a_i | x, t) \propto \vec{m}_i(a_i) \ell_t(x_i | a_i) \overleftarrow{m}_i(a_i)$ and the pairwise marginal on edge $(i-1, i)$ is $\propto \vec{m}_{i-1}(a_{i-1}) \ell_t(x_{i-1} | a_{i-1}) W(a_i | a_{i-1}) \ell_t(x_i | a_i) \overleftarrow{m}_i(a_i)$. In (7) the local likelihood is absorbed *before* propagation, so $\ell_t(x_i | a_i)$ appears exactly once in each marginal; this convention is fixed throughout.

Remark 1 (What “exact” means here). *Three statements must be kept apart. (i) Sum-product on the posterior (6) is exact at the level of functional messages, as a consequence of the graph being a tree. (ii) Messages are functions on $\mathbb{R}$, so an implementation must represent them; we use a truncated uniform grid with composite trapezoidal weights, which defines a finite-state model for which the recursion is exact up to floating-point arithmetic. (iii) That finite-state model approximates the continuous one, with an error that must be measured rather than assumed. We write “exact tree inference” for (i) and “exact for the discretised model” for (ii), and never abbreviate either to “exact”.*




Figure 1: Factor graph of the posterior (6). Latent variables a_i (open) carry binary transition factors along the chain; each observation x_i (shaded) enters through a *unary* factor ℓ_t attached to a single latent variable. Conditioning therefore reweights node potentials without creating edges among the a_i , and the posterior graph remains a chain. Sum-product with the messages (7) is exact on it.

4 Numerical control

Two error sources are separated by construction. *Truncation*: probability mass outside the grid $[-A, A]$ is discarded, monitored by the mass in the boundary cells. [PENDING] The diagnostic is computed inside the recursion but not persisted, so no aggregate figure is quoted. *Quadrature*: the trapezoidal rule is second order in the spacing for integrands with a bounded second derivative. For the Gaussian kernel the integrand is smooth and convergence is fast; **for the Laplace kernel it is not**, since $\varphi(e) \propto e^{-|e|/b}$ has a discontinuous derivative at the origin, and we therefore make no claim of spectral accuracy for that family.

Against the closed-form Gaussian posterior mean, the discretised recursion agrees to 9.2×10^{-15} relative at $K = 401$, $A = 8$. That number measures the discretisation of the Gaussian model only. It is not comparable with the independent Monte Carlo check below, which measures something else: on the Laplace chain, where no closed form exists, an independent importance-sampling estimate agrees with the discretised recursion: the evidence to within five Monte Carlo standard errors at 10^6 samples, and the posterior mean to 10^{-2} relative at 2×10^6 (tests/test_em_bp.py). These are assertions in the test suite rather than recorded measurements, and their tolerances are set by the Monte Carlo error rather than by the grid. Both are reported because they validate different things; neither is the accuracy of the other.

5 Gaussian closure and the LMMSE estimator

A standard approximation propagates only the first two moments of each message. Under the transition $a_i = \rho a_{i-1} + \varepsilon_i$, a message with mean m and variance v transforms as

$$\text{forward: } m \mapsto \rho m, \quad v \mapsto \rho^2 v + q; \quad \text{backward: } m \mapsto m/\rho, \quad v \mapsto (v + q)/\rho^2, \quad (8)$$

the backward map being the forward one inverted because that recursion integrates over the *output* variable (it requires $\rho \neq 0$), and a Gaussian message times a Gaussian likelihood is Gaussian. The propagation therefore uses only the mean and variance of φ and discards everything else about its shape.

Proposition 1. *For the chain (2) with linear transitions, Gaussian unary likelihoods and $\mathbb{E}[a] = 0$, moment-projected sum-product returns the LMMSE estimator of the covariance-matched Gaussian model,*

$$\hat{a}_{\text{LMMSE}} = \text{Cov}(a, x) \text{Cov}(x)^{-1} x = e^{-t} \Sigma_0 (e^{-2t} \Sigma_0 + \Delta_t I)^{-1} x, \quad (9)$$

for every innovation law with mean zero and variance q .

Proof. Both maps in (8) depend on φ only through its mean and variance, so the recursion returns the same function of x for every innovation law with moments $(0, q)$ — in particular the same as on the covariance-matched *Gaussian* chain. On that chain every message is exactly Gaussian, so the moment projection is the identity and the recursion is exact; exact inference returns the posterior mean, which for jointly Gaussian zero-mean (a, x) coincides with the LMMSE estimator by Proposition ?? restricted to linear estimators. $\square$

On the Laplace chain the information-form recursion is verified against the closed-form Gaussian posterior mean to 10^{-12} (`tests/test_gaussian_bp_equivalence.py`). [PENDING] The corresponding check across the full set of innovation families is not currently recorded as an output. The zero-mean hypothesis is not cosmetic: with a prior mean $\mathbb{E}[a] = 1.3$ the estimator (9) departs from the true LMMSE by 0.65 in maximum absolute value, since the general form carries $\mathbb{E}[a] + \text{Cov}(a, x) \text{Cov}(x)^{-1}(x - \mathbb{E}[x])$.

Remark 2. *Proposition 1 is a statement about this closure — single-Gaussian moment projection with linear transitions — and not about Gaussian approximations in general. Within its scope it has a consequence worth stating: what one might call the message-representation error and the model error coincide, because projecting the messages and replacing the model by its covariance-matched Gaussian give the same estimator. The two separate again for richer message families.*

The gap between (9) and exact inference locates where non-Gaussianity matters. It is largest at *small* t and decays monotonically: on a Laplace chain at $\rho = 0.85$, the relative deviation of the closure’s score from the exact one falls from 0.198 at $t = 0.08$ to 9.9×10^{-5} at $t = 2.4$. Large t Gaussianises the posterior, so the shape of φ survives only where the likelihood is weak.

6 Estimating the transition kernel

We now treat W as unknown and observe only noised sequences x^μ , $\mu = 1, \dots, M$, possibly at several noise levels. The initial law $\mu(a_1)$ is held fixed (see §9). Write θ for the kernel parameters and $\mathcal{L}(\theta) = \sum_\mu \log P_{t,\theta}(x^\mu)$.

Parametrisation. We use a mixture-innovation kernel

$$W_\theta(a' | a) = \sum_{c=1}^C \pi_c \mathcal{N}(a' - \rho a; \nu_c, s_c^2), \quad (10)$$

which keeps the linear autoregression and leaves the innovation law essentially free. At $C = 4$ this has ρ , four weights π_c , four locations ν_c and four scales s_c^2 : thirteen numbers subject to $\sum_c \pi_c = 1$, hence **twelve free parameters**. It is never given the density it is asked to recover.

Exact expectation step. Many latent-variable models require an approximate or variational E-step; this one does not, because the posterior is a tree (§3) and sum-product returns exact single-site *and pairwise* marginals. The complete-data log-likelihood is a sum over edges, so the θ -dependent part of the auxiliary function $\mathcal{Q}(\theta | \theta^{(r)})$ — it also contains a site-one term and a likelihood term, both θ -free and therefore additive constants — depends on the data only through

$$\Xi_{kl} = \sum_{\mu=1}^M \sum_{i=2}^n P(a_{i-1} = u_k, a_i = u_l | x^\mu), \quad (11)$$

the expected transition mass between grid points u_k and u_l — the continuum analogue of the expected transition counts of Baum–Welch [Baum et al., 1970]. Sufficiency of (11) in that sense requires a *homogeneous* kernel shared across sites and sequences, a fixed grid, an initial law excluded from θ , and θ -free likelihood factors — the last being what licenses observations at several noise levels to add into one Ξ . It fails if any of these is relaxed. Given those, the maximisation step never revisits the data and costs $O(PK^2)$ for P parameters, independently of M ; observations at different noise levels add into one Ξ . Forming Ξ costs $O(MnK^2)$ per iteration and $O(K^2)$ storage, so “one matrix” is a statement about sufficiency, not about inference being cheap.

Gradient without differentiating the recursion. Fisher’s identity [Fisher, 1925, Cappé et al., 2005] gives

$$\nabla_\theta \mathcal{L}(\theta) = \sum_\mu \mathbb{E}_\theta[\nabla_\theta \log P_\theta(a, x^\mu) | x^\mu] = \langle \Xi(\theta), \nabla_\theta \log W_\theta \rangle, \quad (12)$$

the exact gradient of the *marginal* log-likelihood from one sum-product pass. Both sides are summed over the M observed sequences, matching the definition of Ξ . Sum-product is not a computation graph to backpropagate through; it supplies the conditional expectation. We verify (12) against finite differences of the exact evidence to $\sim 10^{-9}$.

Table 1: Relative denoising error against exact inference under the true kernel, averaged over the noise schedule. Mean over six training seeds, $\pm$ one standard error. The seeds vary initialisation at fixed data, so these intervals measure optimisation variance and not sampling variance.

sequences M	network (25,248 params)	EM-BP (12 free)	ratio
32	0.6070 ± 0.0040	0.0651 ± 0.0031	9.3
128	0.4700 ± 0.0023	0.0333 ± 0.0035	14.1
512	0.2361 ± 0.0016	0.0211 ± 0.0022	11.2
1024	0.1736 ± 0.0009	0.0182 ± 0.0021	9.5
2048	0.1372 ± 0.0010	0.0142 ± 0.0009	9.7
4096	0.1240 ± 0.0003	0.0124 ± 0.0005	10.0

Equation (12) is not EM. It licenses *gradient ascent* on $\mathcal{L}$, which is the procedure described in the guidance we were given, whereas EM maximises $\mathcal{Q}$ at each step. Both are implemented and they do not dominate one another: on the smooth Gaussian kernel gradient ascent reaches EM’s optimum to 2×10^{-9} nats while requiring a tuned step size, and on the Laplace kernel it attains a higher likelihood than EM, because there the exact maximiser of $\mathcal{Q}$ is a lattice artefact of the discretisation. Results below use EM. For the mixture kernel (10) the maximisation step is itself iterative — the complete-data problem carries a second latent variable, the mixture label, and we run four inner conditional-maximisation sweeps — so it is *generalised* EM: each step increases $\mathcal{Q}$ without exactly maximising it, which preserves monotone ascent but not the exact-maximiser property.

Identifiability. Two claims must be separated. Gaussian convolution at finite t is injective on distributions — the characteristic function acquires a factor $e^{-\Delta_t \|\xi\|^2/2}$, which has no zeros, and $e^{-t} > 0$ — so P_t determines P_0 . *Kernel* identifiability does not follow: it additionally requires the initial law to be fixed or separately identifiable, the transition family to be identifiable as a parametric family, and the mixture (10) to be treated modulo label permutation. Two further conditions are easy to overlook. Finite-mixture identifiability holds only on the interior of the parameter set — all $\pi_c > 0$ and the pairs (ν_c, s_c) distinct — which is exactly where an over-specified C need not stay; and W is pinned down only on the support of μ , so a full-support initial law is required. We assume these rather than prove them. We also note that the fitted innovation is not constrained to have zero mean: imposing $\sum_c \pi_c \nu_c = 0$ by reprojection is not a maximisation step and broke monotone ascent in testing, so the condition is monitored as a diagnostic (it lands at the 10^{-3} level) rather than enforced. Proposition 1 therefore applies to the true kernel, not automatically to the estimated one.

What noising does destroy is *information*, and unevenly. Measured per sequence from exact single-chain passes, the Fisher information for the innovation scale falls by a factor 142 between $t = 0.05$ and $t = 1.6$, while that for the correlation falls by 26. The higher-order structure of φ is thus far more fragile under the channel than its second-order structure — a fact we return to in §8.

7 Denoising accuracy

Protocol. Laplace chain, $n = 32$, $\rho = 0.85$. Estimators are scored against exact-BP denoising under the true kernel on held-out sequences, at five noise levels. Neural baselines are trained by denoising score matching in both standard parametrisations (predicting the noise, and predicting the clean signal); we report the better of the two at each noise level, since they win in different regimes and reporting one alone would mislead. [PENDING] Architectural and optimisation details are being written up as an appendix; they are currently read off `src/denoiser.py` $\rightarrow$ `train_dsm_denoiser` and `src/local_head.py` $\rightarrow$ `train_local_head`.

Table 1 shows a ratio between 9 and 14 across seven doublings of the data, with no trend in the budget. The network column takes, at each noise level, the better of the two parametrisations, as the protocol above states; reporting the noise-predicting parametrisation alone — as an earlier version of this table did — would overstate the network’s deficit by 15 to 30%, since predicting the clean signal is the better choice at $t \in \{0.8, 1.6\}$ at every budget. The comparison at equal budget is the meaningful one. We note that the network at $M = 4096$ did not reach the error the estimator attains at $M = 32$; **we did not observe a crossing**, so we do not quantify a data-equivalence factor, and the extrapolation that would be required to do so spans several decades.

Table 2: Generated-sequence statistics at $M = 2048$, Laplace chain, against closed-form targets for $P_{t_{\min}}$. Distributional rows are medians over four seeds; the denoising row is a median over seeds and the five probed noise levels. The denoising column repeats the ranking of §7; the distributional columns do not agree with it.

	EM-BP	local CNN	global MLP	target
denoising MSE against exact	0.00052	0.00525	0.01330	0
innovation excess kurtosis	0.897	1.357	0.206	1.910
$D_{\text{KL}}(\text{true} \parallel \text{generated})$	0.0055	0.0028	0.0165	0
worst covariance-lag error	0.0292	0.0973	0.0518	0

Controls. Capacity is not the explanation: over 24 configurations from 3.2k to 297k parameters the best network reaches 0.208 against the estimator’s 0.022 at $M = 1024$, and error degrades for the largest models. The baseline is competent: trained on a Gaussian chain where the optimal denoiser is linear and closed-form, with 3000 sequences and 12000 steps, it passes the competence threshold asserted in `tests/test_em_bp.py`. The achieved value is not persisted, so it is not quoted.

A locality-respecting baseline. A fully connected network carries no sequential prior. Against a weight-shared window predictor — a one-dimensional convolution — the margin narrows substantially, from roughly 10 to between 2 and 4 over $M \in \{128, 512, 2048\}$. Between 58% and 73% of the fully connected network’s deficit was therefore its architecture rather than the estimator, depending on the budget. The margin against the convolution increases over the three budgets we ran — 2.31 ± 0.39 , 3.36 ± 0.30 , 4.16 ± 0.24 as ratios of seed-averaged errors, or 2.72, 3.78, 5.04 if per-noise-level ratios are formed first. Three points do not establish a trend, and the aggregation choice moves the value, so we report the range. **These particular numbers used a receptive field selected on the evaluation set** and are exploratory pending a rerun with a validation split; the bias favours the baseline, so the margin quoted is if anything conservative.

8 Generation

Denoising accuracy is not what a generative model is judged on. Errors enter the reverse integration at every step and accumulate unevenly across the schedule, so a pointwise ranking need not survive. We therefore integrate the reverse process $dx/d\tau = x + 2S(x, \tau) + \sqrt{2}\eta$ under each estimator (the noise η is standardised throughout, so the $\sqrt{2}$ is explicit on both the forward and the reverse process), with common random numbers across estimators.

Comparing like with like. Integration must stop at $t_{\min} > 0$, since S diverges as $\Delta_t \rightarrow 0$. What each estimator produces is therefore a sample of $P_{t_{\min}}$, not of P_0 , and comparing it against clean data measures the stopping point rather than the score: independent Gaussian noise of variance v on an innovation of variance q multiplies the excess kurtosis by $(q/(q+v))^2$. We instead draw a reference from P_0 and noise it forward to the same $t_{\min}$, so both sides are samples of the same law and the target moments follow in closed form. Run with the *true* kernel, the sampler reproduces that target to within 1.4% across budgets — medians 1.897, 1.883, 1.914, 1.908 at $M = 32, 128, 512, 2048$ against a closed-form 1.9098 — which validates the integrator and the statistics before any estimator is compared. The spread is sampling noise and is not systematic in M .

Table 2 is the paper’s main empirical finding, and it is a dissociation. The structured estimator is an order of magnitude more accurate pointwise and three times better on second-order structure, and yet the convolutional network reproduces the *innovation law* more faithfully, by both kurtosis and divergence. The fully connected network shows the pattern most sharply: over the same range its denoising error improves twelvefold while its generated kurtosis degrades from 0.850 to 0.206. It fits the bulk of the posterior and loses the tail.

A plausible mechanism, not yet established. Both the marginal likelihood that EM maximises and the squared-error denoising loss are dominated by the bulk of the distribution. The tail of φ carries little probability mass and therefore little weight in either objective, while contributing disproportionately to the shape statistics that reverse diffusion exposes. This is the same asymmetry

measured in §6: the channel destroys innovation-shape information 142-fold against 26-fold for the correlation, and the estimator sees only noised data. Whether the deficit is a capacity limit of the parametrisation (10) or an information limit of the observation model is decided by varying C , which is in progress and not reported here.

9 Limitations and scope

The model assumes exact first-order Markov structure; under a non-Markov prior the estimator is misspecified in a way more data does not repair. The transition family (10) may additionally be misspecified, and §8 suggests it is, for the tails. The initial law is fixed rather than estimated. Sum-product is exact on the posterior chain but the implementation is a discretisation whose error, for the Laplace family, is second order rather than spectral.

Inference costs $O(nK^2)$ per sequence per evaluation against a few matrix products for a network, measured here at two to three hundred times slower; since reverse diffusion calls the denoiser at every step, this is the practical weak point. Training is also slower, and the estimator’s cost grows linearly in M while the network’s is fixed by its step count, so the advantage is in data and not in wall-clock. The low parameter count is not free: it is bought with a structural assumption, and the advantage narrows as the parametric dimension grows — on a discrete-alphabet variant it falls from 15 to 4.4 as the estimated dimension goes from 6 to 132 parameters.

Error bars over training seeds measure optimisation variance at fixed data. The generative comparison is complete for one innovation family. Comparisons are against the two architectures described and do not license a claim about neural denoisers in general.

10 Conclusion

When a sequential distribution is known to have low-dimensional Markov structure, learning a diffusion score reduces to latent-variable inference and estimation of a compact transition model. Sum-product on the noised chain supplies an exact non-Gaussian benchmark, and estimating the kernel from noised data alone gives a denoiser that is substantially more accurate at small sample sizes than trained networks, at the cost of stronger assumptions and much more expensive inference. The comparison after reverse diffusion does not reproduce that ranking, and we regard understanding the dissociation — rather than the pointwise margin itself — as the useful direction.

References

- Leonard E. Baum, Ted Petrie, George Soules, and Norman Weiss. A maximization technique occurring in the statistical analysis of probabilistic functions of markov chains. *The Annals of Mathematical Statistics*, 41(1):164–171, 1970.
- Giulio Biroli, Tony Bonnaire, Valentin de Bortoli, and Marc Mézard. Dynamical regimes of diffusion models. *Nature Communications*, 15:9957, 2024.
- Olivier Cappé, Eric Moulines, and Tobias Rydén. *Inference in Hidden Markov Models*. Springer, 2005.
- Ronald A. Fisher. Theory of statistical estimation. *Mathematical Proceedings of the Cambridge Philosophical Society*, 22(5):700–725, 1925.
- Jérôme Garnier-Brun, Marc Mézard, Emanuele Moscato, and Luca Saglietti. How transformers learn structured data: insights from hierarchical filtering, 2024.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *Advances in Neural Information Processing Systems*, 2020.
- Mason Kamb and Surya Ganguli. An analytic theory of creativity in convolutional diffusion models, 2024.
- Song Mei and Yuchen Wu. U-nets as belief propagation, 2024. See also related work on hierarchical denoising.

- Marc Mézard. Generative diffusion: updated notes, 2025. Lecture notes, Cargèse summer school.
- Marc Mézard and Andrea Montanari. *Information, Physics, and Computation*. Oxford University Press, 2009.
- Judea Pearl. *Probabilistic Reasoning in Intelligent Systems*. Morgan Kaufmann, 1988.
- Antonio Sclocchi, Alessandro Favero, and Matthieu Wyart. A phase transition in diffusion models reveals the hierarchical nature of data. *Proceedings of the National Academy of Sciences*, 122(1): e2408799121, 2025.
- Jascha Sohl-Dickstein, Eric Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. In *International Conference on Machine Learning*, 2015.
- Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations*, 2021.