

BP-Informed Diffusion Scores for Approximately Markovian Sequences

Calibrated grid BP, the message/model equivalence for single-Gaussian projections, rank-one non-Markov corrections, hybrid residual learning, and reverse dynamics

Thesis project continuation — second iteration

July 2026

Abstract

This report continues the study of belief-propagation (BP) computation of diffusion scores for sequence data with (approximately) Markov structure. Building on the previous note, we (i) audit and exactly reproduce the earlier package, identifying a numerical boundary-collapse artifact in its grid-projected Gaussian BP that contaminated all weak-noise results; (ii) prove that single-Gaussian moment-matched (assumed-density) BP on any *linear-transition* chain with Gaussian observation factors is identical to exact Gaussian BP under the covariance-matched Gaussian model, so that “Gaussian message error” and “Gaussian model error” coincide at this level; (iii) calibrate grid BP against the exactly solvable Gaussian AR(1) case, finding *spectral* accuracy with an abrupt small- t failure driven by likelihood aliasing; (iv) quantify the corrected Gaussian-baseline error across heavy-tailed and bimodal innovation families; (v) show that for Gaussian AR(1)-plus-global-latent priors the non-Markov score correction is *exactly rank one* at every diffusion time, and that a neural network learning only this residual on top of a frozen Markov-BP score is an order of magnitude more sample- and parameter-efficient than a direct score network; and (vi) demonstrate in reverse-diffusion simulations that the Gaussian score reproduces second-order statistics while completely washing out heavy-tailed innovations, even though its trajectory-level effect is dynamically stable in the L^2 sense. Throughout we separate exact statements from numerical approximations.

Contents

1	Setting and exact results	3
1.1	Model and forward convention	3
1.2	Posterior chain, BP exactness, score identity	3
1.3	The master error identity	3
2	Audit of the previous package (Layer 1)	4
3	Gaussian closure, and what a single Gaussian can ever measure	5
4	Layer 2: calibrating grid BP on the exactly solvable case	6
5	Layer 3: the corrected Gaussian baseline on non-Gaussian chains	8
5.1	Laplace chain, per-trial analysis (exp. 02)	8
5.2	Innovation sweep (exp. 03)	9

6	Layer 4: approximate Markovianity	9
6.1	Model A: AR(1) plus global latent — the residual is rank one	9
6.2	Model B: weak long-range precision coupling	10
6.3	Hybrid learning: BP baseline plus neural residual	11
7	Reverse-diffusion dynamics (exp. 05)	12
7.1	Does the Gaussian score wash out non-Gaussian structure?	12
7.2	Pointwise error vs dynamical error	12
7.3	Model A: recovering the global mode	13
8	Interpretation, limitations, and next steps	13
A	Package layout and reproduction	14
B	Relation to the previous report	14

1 Setting and exact results

1.1 Model and forward convention

Clean data $a_{1:n} \in \mathbb{R}^n$ follow a Markov chain prior

$$p_0(a_{1:n}) = \mu(a_1) \prod_{i=2}^n K_i(a_i | a_{i-1}). \quad (1)$$

The forward noising process is the coordinatewise Ornstein–Uhlenbeck (variance-preserving) diffusion

$$dX_t = -X_t dt + \sqrt{2} dW_t, \quad X_0 = a, \quad (2)$$

whose transition kernel gives, at time t ,

$$x_i = \alpha_t a_i + \sqrt{\Delta_t} z_i, \quad \alpha_t = e^{-t}, \quad \Delta_t = 1 - e^{-2t}, \quad z_i \sim \mathcal{N}(0, 1) \text{ i.i.d.} \quad (3)$$

For unit-variance data, $\text{Var}(x_i) = \alpha_t^2 + \Delta_t = 1$ at every t ; all priors in this report are normalized to unit marginal variance so that $p_T \rightarrow \mathcal{N}(0, I)$. The sitewise likelihood is the exact Gaussian factor $\ell_i(a_i; x_i) = \mathcal{N}(x_i; \alpha_t a_i, \Delta_t)$.

1.2 Posterior chain, BP exactness, score identity

Because the prior factorizes over adjacent pairs and the likelihood factorizes over sites, the posterior is again a chain,

$$p_t(a | x) \propto \mu(a_1) \prod_{i=2}^n K_i(a_i | a_{i-1}) \prod_{i=1}^n \ell_i(a_i; x_i), \quad (4)$$

a tree-structured factor graph on which sum-product BP computes exact marginals $b_i(a_i) = p_t(a_i | x)$ after one forward and one backward sweep (full derivations in the previous report; they were checked line by line during the audit and are correct). Combined with the denoising (Tweedie) identity for the OU kernel,

$$\boxed{s_i(x, t) = \partial_{x_i} \log p_t(x) = -\frac{x_i - \alpha_t m_i(x, t)}{\Delta_t}, \quad m_i(x, t) = \mathbb{E}[a_i | x_{1:n}],} \quad (5)$$

this yields the exact theorem: *for Markov-chain priors under coordinatewise OU noising, BP computes the exact diffusion score*, provided messages are represented exactly. The messages, however, are continuous functions; outside conjugate families they admit no finite parametrization. Everything that follows concerns the quantitative cost of representing them approximately, and what survives when the data are only *approximately* Markov.

1.3 The master error identity

If $\hat{m}$ is any approximate posterior mean and $\hat{s} = -(x - \alpha_t \hat{m})/\Delta_t$, then for any reference pair $(m_{\text{ref}}, s_{\text{ref}})$ produced the same way,

$$\boxed{\hat{s} - s_{\text{ref}} = \frac{\alpha_t}{\Delta_t} (\hat{m} - m_{\text{ref}}), \quad \frac{\alpha_t}{\Delta_t} \sim \frac{1}{2t} \quad (t \downarrow 0).} \quad (6)$$

All score errors are amplified posterior-mean errors, with amplification diverging at small noise. Every experiment in this package logs the numerical residual of (6) (column `identity_residual`);

it is at machine precision (10^{-15} – 10^{-12}) throughout, as it must be, since it is algebra, not approximation. (When the compared scores agree to machine precision the *relative* residual is 0/0 and meaningless; the tables flag this case.)

Two further decompositions organize the study. Against an exact score,

$$\hat{s} - s_{\text{true}} = \underbrace{(\hat{s} - s_{\text{grid}})}_{\text{message/model approx.}} + \underbrace{(s_{\text{grid}} - s_{\text{true}})}_{\text{grid discretization}}, \quad (7)$$

and for approximately Markov data,

$$s_{\text{true}} = s_{\text{Markov}} + s_{\text{residual}}, \quad s_{\text{hybrid}, \theta} = s_{\text{BP}, K_{\varphi}} + r_{\theta}. \quad (8)$$

2 Audit of the previous package (Layer 1)

The previous package was rerun under its documented configuration (`-mode report -n-trials 12`, seed 11). All three CSV outputs reproduce *bit-for-bit up to machine precision* (largest deviation 2.1×10^{-14} , attributable to numpy/BLAS version differences). The mathematical derivations match the code. Two substantive findings change the interpretation of its results.

F1: boundary collapse of grid-projected Gaussian messages. The legacy `gaussian_projected_bp` evaluates each Gaussian message on the grid, applies the exact grid update, and moment-matches the result. At weakly informative times the outgoing message function is a near-flat ramp whose maximizer lies *outside* $[-A, A]$; moment-matching the truncated ramp yields a mean near the grid boundary and a spuriously small variance, and the next update pushes it further edge-ward (the backward integrand in a peaks at a'/ρ , outside the grid) — a positive feedback loop. Diagnosed example (Gaussian prior, $t = 1.8$): backward messages collapse to $R_{\text{mean}} \approx -7.7$, $R_{\text{var}} \approx 0.08$ on a $[-8, 8]$ grid, producing a posterior-mean error of 6.9 at one site. This artifact—not Gaussian closure—explains the anomalous large- t rows of the previous package (posterior-mean MSE 2.2 ± 7.5 at $t = 1.3$ on the Laplace chain, and MSE up to 5.0 in the *Gaussian* sanity check where projection should be exact). Figure 1 quantifies the artifact against the corrected implementation: it exceeds 6 in posterior-mean sup-norm for $t \geq 0.9$, precisely where the true closure error is below 10^{-2} .




Figure 1: Audit finding F1. Sup-norm difference between the legacy grid-projected Gaussian BP posterior mean and the analytic information-form result on identical data (Laplace chain, $\rho = 0.85$). The two implement the same projection map, so the difference is purely numerical. The large- t explosion is the boundary-collapse artifact.

F2: message error = model error for single-Gaussian projections. See Proposition 1 below: what the previous package measured as “Gaussian *message* approximation error” on the Laplace chain is mathematically the error of the covariance-matched Gaussian *model*. The numbers were right; the interpretation needs sharpening. The distinction between message-level and model-level approximation only becomes real for richer message families (mixtures, particles) or nonlinear transitions.

Minor findings: grid sizes were compared on different random draws (fixed by common random numbers keyed on (ρ, t, trial)); linear-domain likelihood rows can underflow entirely for small t (fixed by row-shifted log-likelihoods, exact because per-site constants cancel); ad-hoc variance floors hid resolution failures (replaced by an explicit resolution check). The corrected package returns, in addition, the exact model evidence $\log p_t(x)$ from the BP normalization constants; it matches the closed-form Gaussian marginal likelihood in the calibration case (unit test).

3 Gaussian closure, and what a single Gaussian can ever measure

In the Gaussian AR(1) case messages close in information form (λ, h) : multiplication by the likelihood adds $(\alpha_t^2/\Delta_t, \alpha_t x_i/\Delta_t)$, and the linear transition maps moments $(m, v) \mapsto (\rho m, \rho^2 v + q)$. The corrected implementation (`gaussian_chain_bp`) performs exactly these updates analytically; a flat message is $\lambda = 0$ exactly, so no truncated moment-matching ever occurs.

Proposition 1 (Single-Gaussian ADF-BP $\equiv$ matched Gaussian model BP). *Let the clean chain have linear transitions $a_i = \rho a_{i-1} + \varepsilon_i$ with $\mathbb{E}\varepsilon_i = 0$, $\text{Var}(\varepsilon_i) = q$, arbitrary innovation law, Gaussian initial density, and Gaussian sitewise likelihoods (3). Then assumed-density BP that projects every message to a single Gaussian by moment matching produces exactly the messages, beliefs, posterior means, and score of exact BP under the Gaussian AR(1) prior with the same (ρ, q) .*

Proof. Forward pass: if the incoming message L_i is Gaussian, the tilted product $L_i \cdot \ell_i$ is a product of Gaussians, hence exactly Gaussian $\mathcal{N}(\tilde{m}, \tilde{v})$ — projection is the identity here. The outgoing message is the law of $\rho\xi + \varepsilon$ with $\xi \sim \mathcal{N}(\tilde{m}, \tilde{v})$; whatever the innovation law, its first two moments are $(\rho\tilde{m}, \rho^2\tilde{v} + q)$, so the moment-matched Gaussian is $\mathcal{N}(\rho\tilde{m}, \rho^2\tilde{v} + q)$ — identical to the Gaussian-innovation update. Backward pass: the outgoing function is $a \mapsto \mathbb{E}_\varepsilon[\tilde{g}(\rho a + \varepsilon)] = (\tilde{g} * \check{p}_\varepsilon)(\rho a)$ with $\tilde{g} = \mathcal{N}(\tilde{m}, \tilde{v})$; the convolution has mean $\tilde{m}$ and variance $\tilde{v} + q$ regardless of the innovation shape, so after the affine change of variable the moment-matched message is $\mathcal{N}(\tilde{m}/\rho, (\tilde{v} + q)/\rho^2)$, again the Gaussian-model update. By induction from the Gaussian initializations ($L_1 = \mu$ Gaussian, R_n flat), all message parameters coincide; beliefs are products of Gaussians and coincide as well. $\square$

Remark 1. *The projection discards the innovation law beyond its variance at every step, so nothing non-Gaussian ever enters. Consequently: (i) the honest implementation of the single-Gaussian baseline is the analytic information-form filter (no grid); (ii) measuring “message error” distinct from “model error” requires message families that can carry non-Gaussian information (Gaussian mixtures being the natural next step — with mixture incoming messages the tilted products are mixtures and multimodality survives the transition step); (iii) the proposition fails for nonlinear transitions, where the tilted forward moments genuinely depend on the innovation shape. This is verified numerically: the legacy grid projection agrees with the analytic filter to 5×10^{-3} sup-norm at moderate t on the Laplace chain (unit test), the residual being grid quadrature error.*

4 Layer 2: calibrating grid BP on the exactly solvable case

Grid BP represents messages by values on M equispaced points in $[-A, A]$ with trapezoidal quadrature. Experiment 01 sweeps $M \in \{101, 201, 401, 801\}$, $A \in \{4, \dots, 8\}$, $t \in [0.005, 2.4]$, $\rho \in \{0.2, 0.5, 0.8, 0.95\}$ against the exact Gaussian score, with identical random data across all grid configurations.




Figure 2: Relative score error of grid BP vs the exact precision-matrix score ($\rho = 0.95$, $A = 8$). The error sits at the 10^{-15} floor except where the likelihood width $\sqrt{\Delta_t}/\alpha_t \approx \sqrt{2t}$ approaches the grid step: failure is abrupt, not gradual.

Spectral accuracy. The trapezoid rule on analytic, rapidly decaying integrands converges exponentially in points-per-width (aliasing error $\sim e^{-2\pi^2\sigma^2/dx^2}$), not at the naive $O(dx^2)$ rate. Hence: at $M = 101$, $A = 8$ ($dx = 0.16$) the error is 10^{-15} down to $t = 0.02$, degrades to 10^{-4} at $t = 0.01$, and fails (1.7×10^{-2}) at $t = 0.005$; at $M = 201$ it is at the floor for all tested t . Tail truncation is visible only at $A = 4$ (error 10^{-6} – 10^{-4} depending on ρ ; at $A = 5$, $\leq 7 \times 10^{-6}$ for Gaussian tails). See Figures 3.




Figure 3: Left: error vs M at $A = 8$. Right: error vs A at fixed $dx = 0.02$ (truncation isolated).




Figure 4: (M, A) heatmaps of $\log_{10}$ relative score error at $\rho = 0.95$; $t = 0.005$ (left) shows the aliasing-driven failure region, $t = 0.4$ (right) shows the truncation-only regime.

Recommended defaults. The Gaussian calibration alone would admit $M = 101$, $A = 5$; because heavy-tailed innovations put more posterior mass in the tails, we adopt $M = 401$, $A = 8$ as the working reference for all non-Gaussian experiments (error floor in calibration, cost $O(nM^2)$ still trivial) and $M = 801$ for convergence checks. The binding small- t constraint is $\sqrt{2}t \gtrsim 3dx$, checked automatically (`resolution_ok`).

5 Layer 3: the corrected Gaussian baseline on non-Gaussian chains

5.1 Laplace chain, per-trial analysis (exp. 02)

On the Laplace chain ($\rho = 0.85$, $n = 40$, 60 trials per t , reference $M = 801$, $A = 8$), the corrected baseline gives a clean monotone picture (Figure 5): median relative score error 0.39 at $t = 0.02$, 0.17 at $t = 0.08$, 3.5×10^{-2} at $t = 0.4$, 7.6×10^{-5} at $t = 2.4$. The failure probability $P(\text{err} > 0.1)$ is 1 for $t \leq 0.08$ and 0 for $t \geq 0.9$. The grid self-convergence error ($M = 401$ vs 801) remains 10^4 – 10^7 times smaller than the closure error at every t : the reference is trustworthy where it is used. The old package’s large- t anomalies are absent — they were artifact F1.




Figure 5: Gaussian closure error (median, 10–90% band) vs the grid error budget on identical data. Laplace chain, $\rho = 0.85$.




Figure 6: Worst trial across the sweep: the reference belief is sharply non-Gaussian and the single-Gaussian belief commits to a compromise; by (6) the mean shift is amplified by α_t/Δ_t .

5.2 Innovation sweep (exp. 03)

Six families with *identical* covariance $\rho^{|i-j|}$ and varying innovation shape: Laplace (excess kurtosis +3), Student- t with $\nu = 5, 8$ (+6, +1.5), symmetric Gaussian mixtures with separation $\kappa = 0.3, 0.6, 0.9$ (−0.18, −0.72, −1.62; bimodal). Findings (Figure 7, full table in `outputs/exp_03.../SUMMARY_TABLE...`).

- **Noise level dominates:** every family’s error decays to $\sim 10^{-3}$ or below by $t = 2.4$; at high noise the Gaussian likelihood dominates the posterior and closure is nearly free. All meaningful error lives at $t \lesssim 0.4$.
- **Both signs of non-Gaussianity hurt, bimodality hurts most:** at $t = 0.05$, $\rho = 0.5$: $\kappa = 0.9$ mixture 0.94, Laplace 0.52, $t_{\nu=5}$ 0.32, $\kappa = 0.3$ mixture 0.07. Strongly bimodal innovations are the worst case at small t , matching the mechanism of Figure 6.
- **Stronger correlation helps, not hurts:** at fixed t the error *decreases* in ρ (Laplace at $t = 0.05$: 0.43, 0.27, 0.13 for $\rho = 0.5, 0.85, 0.95$). Informative neighbours act like extra Gaussian observations and Gaussianize the local tilted posterior. The naive intuition that message errors “propagate farther” at large ρ is not what happens on chains.




Figure 7: Left: median relative score error vs t by family ($\rho = 0.85$). Right: error vs innovation excess kurtosis (negative = bimodal, positive = heavy-tailed) at three times.

By Proposition 1 these numbers are simultaneously the single-Gaussian *message* error and the matched Gaussian *model* error. Separating the two requires mixture messages; this is the clearest next algorithmic step and is left for the next iteration.

6 Layer 4: approximate Markovianity

Here the prior is *not* a chain, but Gaussian, so every score is exact linear algebra and the pure “Markov approximation error” is isolated with no grid and no Monte Carlo in the operator analysis.

6.1 Model A: AR(1) plus global latent — the residual is rank one

Let $a = (y + \beta g)/\sqrt{1 + \beta^2}$ with y a Gaussian AR(1)(ρ) chain and $g \sim \mathcal{N}(0, 1)$ global, so $\Sigma_0 = (\Sigma_y + \beta^2 \mathbf{1}\mathbf{1}^\top)/(1 + \beta^2)$.

Proposition 2 (Rank-one non-Markov correction). *Let $S_t = \alpha_t^2 \Sigma_{\text{chain}} + \Delta_t I$ be the noisy covariance of any chain prior and let the true noisy covariance be $\Sigma_t = S_t + c \mathbf{1}\mathbf{1}^\top$ with $c = \alpha_t^2 \beta^2 / (1 + \beta^2)$. Then by the Sherman–Morrison identity*

$$s_{\text{true}}(x) - s_{\text{chain}}(x) = \frac{c(\mathbf{1}^\top S_t^{-1} x)}{1 + c \mathbf{1}^\top S_t^{-1} \mathbf{1}} S_t^{-1} \mathbf{1}, \quad (9)$$

i.e. the non-Markov score correction is exactly rank one at every t : a fixed smooth direction $S_t^{-1} \mathbf{1}$ with an x -linear scalar coefficient.

Numerically (exp. 04): the residual operator $D_t = \Sigma_{\text{chain},t}^{-1} - \Sigma_t^{-1}$ of the naive chain has $\sigma_2/\sigma_1 \sim 10^{-12}$ (exact rank one), and chain-BP + correction (9) agrees with the exact score to 5×10^{-14} — this *is* augmented BP with the global latent integrated out in closed form. The best-matched chain (Chow–Liu, adjacent pairwise marginals matched) has smaller residual *norm* (e.g. 0.066 vs 0.086 median relative residual at $\beta = 1$, $t = 0.2$) but spreads it over many directions: structure and norm trade off. Residual magnitudes are graceful in β : median relative score residual at $t = 0.2$ of 0.002/0.011/0.03/0.07 for $\beta = 0.1/0.25/0.5/1.0$ (Chow–Liu).




Figure 8: Model A. Left: relative score residual of the best chain approximation. Right: σ_2/σ_1 of the residual operator — the naive chain’s correction is exactly rank one (dashed, at numerical zero); Chow–Liu trades structure for norm.

6.2 Model B: weak long-range precision coupling

With $Q = Q_{\text{AR1}} + \gamma Q_{\text{long}}$ (exponentially decaying long-range couplings, diagonally compensated, covariance renormalized to unit diagonal), the Chow–Liu chain residual grows like 0.084/0.13/0.17/0.23 (median, $t = 0.2$) for $\gamma = 0.05/0.1/0.2/0.4$, with effective rank ≈ 30 : a genuinely high-rank correction. Low-rank latent structure (Model A) is the favourable regime for Markov-plus-residual decompositions; generic long-range couplings are not, although the residual *norm* remains modest for weak γ .




Figure 9: Model B: residual norm (left) and effective rank of the residual operator (right).

6.3 Hybrid learning: BP baseline plus neural residual

On Model A ($\beta = 0.5$, $\rho = 0.8$, $n = 40$), supervised regression on exact-score targets, times log-uniform in $[0.02, 3]$: a direct MLP $s_\theta(x, t)$ versus the frozen Chow–Liu Markov-BP score plus an identical-architecture MLP learning only the residual (8). Test relative score error on 2048 held-out samples:

N_{train}	64	256	1024	4096
direct MLP	0.88	0.75	0.29	0.20
hybrid (BP + residual MLP)	0.23	0.060	0.015	0.014
Markov BP alone (no learning)		0.031		

The hybrid improves on the free Markov baseline from 256 samples on and is $\sim 20\times$ better than the direct network at every size; the direct network does not reach the *no-learning* baseline even at 4096 samples. Parameter efficiency is equally lopsided: a width-16 residual net (1.7k parameters, error 0.017) outperforms a width-256 direct net (87k parameters, error 0.33).




Figure 10: Sample (left) and parameter (right) efficiency of direct vs hybrid score learning on Model A.

Remark 2 (Honest caveats). *The Markov model is known here, not estimated (K_φ given, not learned); the comparison isolates the value of the structural baseline, not of a full estimation*

pipeline. Targets are exact scores (supervised), not denoising score matching; the sample-efficiency gap could change under DSM noise. And Model A is the friendly case: its residual is exactly rank one and x -linear — which is precisely the point: when approximate Markovianity holds with low-rank corrections, the residual is a much easier learning problem than the full score.

7 Reverse-diffusion dynamics (exp. 05)

Convention: reverse SDE $x \leftarrow x + h(x + 2s(x, t)) + \sqrt{2h}\xi$ integrated on a geometric grid from $T = 3$ to $t_{\min} = 0.01$ (80 steps, 200 trajectories); probability-flow ODE $\dot{x} = -x - s$ with Heun steps for reconstruction. Paired comparisons share initial noise and Brownian increments.

7.1 Does the Gaussian score wash out non-Gaussian structure?

Generating from the Laplace chain ($\rho = 0.85$) with the grid-BP score versus the Gaussian-BP score:

	innovation exc. kurtosis	marginal var	lag-1 corr.
clean prior samples	2.72	0.94	0.837
reverse, grid-BP score	2.86	1.06	0.828
reverse, Gaussian-BP score	0.12	1.05	0.832

Both samplers reproduce the second-order statistics; only the grid-BP score recovers the heavy-tailed innovations. This is the dynamical face of Proposition 1: the Gaussian score *is* the score of a Gaussian model, so its reverse diffusion generates (asymptotically) Gaussian samples — the non-Gaussianity is not distorted but *erased*.




Figure 11: Left: innovation kurtosis of reverse samples. Right: pointwise score deviation on the reference path vs cumulative common-noise trajectory divergence.

7.2 Pointwise error vs dynamical error

Along the reverse trajectory, the pointwise Gaussian-vs-grid score deviation is $\sim 10^{-5}$ at $t = 3$, crosses 10^{-1} near $t \approx 0.15$, and reaches 0.49 at $t_{\min}$; yet the cumulative trajectory divergence under common noise saturates around 0.11. Two conclusions. First, the dynamics are *stable*: score error is not amplified along the flow (the late-time contraction of the reverse SDE limits accumulation; errors arrive mostly after the trajectory’s large-scale structure is frozen). Second, stability in L^2 is deceptive: the modest terminal displacement is exactly what carries the higher-moment structure,

as the kurtosis table shows. Small pointwise score error at moderate t , plus stability, is *not* sufficient for distributional correctness in the tails.

Reconstruction from $t = 1$ (probability-flow ODE) shows the two scores recovering correlation with the clean sample nearly identically (second-order task, consistent with the above); see `setup2_correlation_recovery.png`.

7.3 Model A: recovering the global mode

Generating from Model A ($\beta = 0.5$) with four scores: exact, Chow–Liu Markov, naive Markov, and hybrid (chain + rank-one). Top eigenvalue of the terminal sample covariance (true value 12.7; finite-batch/discretization biases affect all samplers equally — exact-score value 10.5): exact 10.5, hybrid 10.5 (bitwise-identical trajectories, as (9) is exact), Chow–Liu 8.1, naive 5.6. A Markov score cannot generate the global collective mode; the rank-one correction restores it exactly.




Figure 12: Model A: top eigenvalue of the sample covariance of reverse-generated samples.

8 Interpretation, limitations, and next steps

What is exactly established. (i) For Markov-chain priors under coordinatewise OU noising, BP computes the exact score; the entire computational question is message representation. (ii) Single-Gaussian projected BP on linear chains equals matched-Gaussian-model BP (Prop. 1); its failure modes are therefore model failures: erased tails and modes, worst at small t by (6). (iii) For chain-plus-low-rank-latent Gaussian priors, the non-Markov score correction is exactly low-rank (Prop. 2); augmented BP is closed-form.

What is numerically established. Grid BP is spectrally accurate and abruptly fails only below a predictable t ; with $M = 401$, $A = 8$ it is a sound reference for every experiment here. The corrected Gaussian baseline degrades smoothly and predictably in $(t, \text{kurtosis}, \rho)$. Hybrid BP-plus-residual learning dominates direct score learning by an order of magnitude on a model with genuine low-rank non-Markovianity. Reverse dynamics under approximate scores are L^2 -stable but distributionally wrong precisely in the structure the approximation discards.

What this does not show. BP does not solve general diffusion models: real priors are unknown, non-chain, and high-dimensional; our non-Gaussian ground truths are numerical (grid), not analytic; the hybrid experiments assume the Markov model known and use supervised score targets. Claims are limited to: *if the clean distribution is (approximately) a Markov chain, BP-based scores are exact (resp. structured baselines), and the non-Markov residual can be dramatically simpler than the full score — in the low-rank-latent case, exactly rank one.*

Next steps, in order of leverage.

1. **Gaussian-mixture messages:** the first family where message error $\neq$ model error; measure how many components close the small- t gap on the bimodal chains (the worst case found here).
2. **Estimated K_φ :** learn the transition from data (parametric or nonparametric), quantify how transition-estimation error propagates through (6) to the score, completing the hybrid pipeline $s_{\text{BP}, K_\varphi} + r_\theta$ end to end.
3. **Denoising-score-matching training** of the residual (no exact targets), to test whether the sample-efficiency gap survives DSM variance.
4. **Model C** (non-Gaussian chain + global latent): Rao–Blackwellized BP with a quadrature over the latent times grid-BP chains — machinery already in place; the reference is small- n quadrature or particle methods.
5. **Nonlinear transitions**, where Proposition 1 fails and moment-matched messages genuinely differ from any Gaussian model — the sharpest place to study message-level approximation per se.

A Package layout and reproduction

Path	Content
<code>src/</code>	library: priors, noising, grid BP (+ evidence), information-form Gaussian BP, exact scores, Markov approximations, numpy MLP, samplers
<code>experiments/exp_01..05</code>	all experiments; <code>-quick</code> smoke mode; CSV/JSON/PNG outputs
<code>outputs/</code>	committed results of the full runs used in this report
<code>notebooks/01..04</code>	executed analysis notebooks
<code>tests/</code>	14 pytest checks (exactness, identities, equivalence, convergence)
<code>audit/AUDIT_NOTE.md</code>	full Layer-1 audit

Determinism: all randomness flows through `src.utils.rng_for` (seeded by experiment/setting keys), so method comparisons are paired by construction. Provenance (package versions, timestamps) is written into every `params.json`.

B Relation to the previous report

The general theorem, its proof, and the Gaussian-closure algebra are unchanged from *From Gaussian AR(1) Belief Propagation to General Markov-Chain Diffusion Scores* (July 2026) and are not

repeated in full here. The Laplace-chain numbers of that report remain valid at $t \leq 0.9$ and are superseded at $t \geq 1.3$ by the corrected baseline (artifact F1). Its interpretation of “Gaussian message error” should be read as “matched Gaussian model error” per Proposition 1.

References

- [1] G. Mantovani, *From Gaussian AR(1) Belief Propagation to General Markov-Chain Diffusion Scores*, internal project note, July 2026.
- [2] M. Mézard and A. Montanari, *Information, Physics, and Computation*, Oxford University Press, 2009.
- [3] B. D. O. Anderson, *Reverse-time diffusion equation models*, Stochastic Processes and their Applications, 1982.
- [4] T. Minka, *Expectation Propagation for approximate Bayesian inference*, UAI 2001. (Assumed-density filtering and moment projection.)
- [5] G. Biroli and M. Mézard, *Generative diffusion in very large dimensions*, arXiv:2306.03518, 2023.
- [6] G. Biroli, T. Bonnaire, V. de Bortoli, M. Mézard, *Dynamical Regimes of Diffusion Models*, arXiv:2402.18491, 2024.
- [7] S. Mei, *U-Nets as Belief Propagation*, arXiv:2404.18444, 2024.